

**MIS 6346.503**

4/21/2023

# Used Cars Project Report

Big Data Final Project

Mahathi Rallapalli - MXR210124

Shivani Kothapally - SXX220116

Sai kruthik reddy paduru - sxp210343

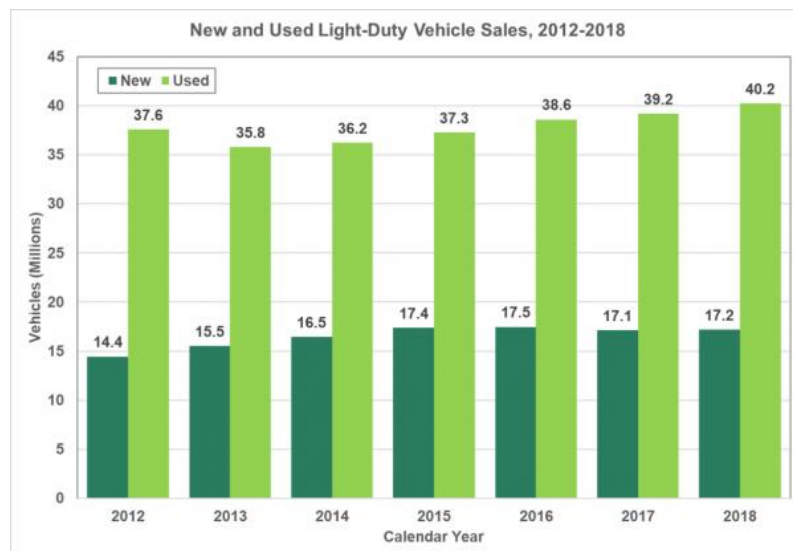
Sindhuja Sai batchu - Sxb220095

Prof. Antonio Paes

## Introduction:

According to recent survey, nearly 70 percent of Americans would be likely to consider a used vehicle for their next auto purchase. There is substantial reasoning for it-

- **Cost:** The lower price of used automobiles as opposed to new cars is one of the most significant factors influencing their choice for them. Once it is driven off the lot, a new car can lose a lot of value, making it a considerably more expensive buy than a used car.
- **Depreciation:** In the first few years of ownership, new cars also experience rapid depreciation, which lowers their resale value. Buyers can escape the first depreciating hit and even make long-term financial savings by choosing a used automobile.
- **Better value:** Compared to new cars, used cars frequently represent a better value. Customers may be able to buy a used automobile with more amenities, greater quality, or higher efficiency for the same price as a new car.
- **Availability:** Thanks to improvements in manufacturing and technology, automobiles are more durable and dependable than before. It is now simpler for customers to select a used car that suits their needs because there are more used cars accessible than ever before.
- **Environmental Concerns:** Due to their negative effects on the environment, some people are also preferring old cars. While buying a used automobile can help lessen the demand for new cars and the resulting environmental impact, producing a new car can often need many resources and energy.



## Problem Statement:

Data analysis and Business answers to a big dataset of used cars in US and Canada.

The used cars dataset contains a massive amount of data that can help businesses in the automotive industry to answer key business questions. Some of the possible business questions that can be addressed using this dataset are:

- Which features have the highest correlation with the price of a car?
- What could be the predicted price and range of popular cars?
- How do different brands and models of used cars compare in terms of price, mileage, and other specifications?

There are several reasons why conducting data analysis on the used cars in the US and Canada dataset could be beneficial:

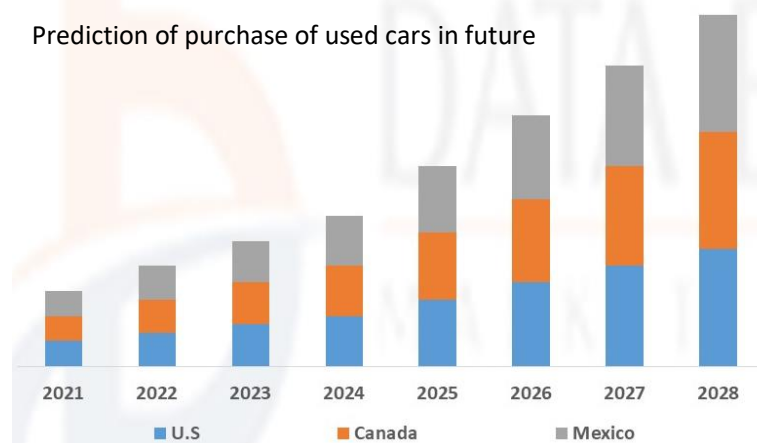
*Market insights:* By analyzing the dataset, researchers can gain insights into the used car market in the US and Canada. This can include identifying trends in prices, sales volume, and popular models, as well as understanding how various factors such as age, mileage, and condition affect pricing.

*Customer behavior:* Analysis of the dataset can also provide insights into customer behavior when purchasing used cars. This can include identifying the most important factors that buyers consider when making a purchase, such as price, mileage, and vehicle features.

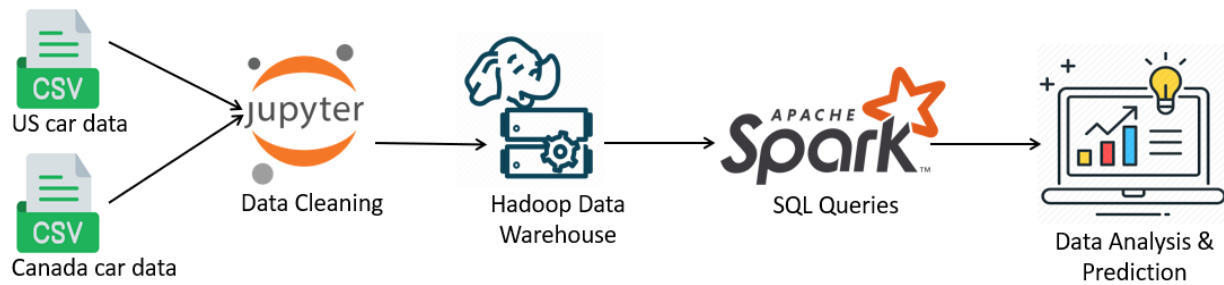
*Business strategy:* The data analysis can be used to help businesses develop strategies for pricing and marketing their used cars. This includes identifying which models and features are in demand, what pricing strategies are most effective, and how to target specific customer segments.

*Predictive modeling:* The data can also be used to develop predictive models that can forecast future trends in the used car market. This can help businesses and investors make informed decisions about buying and selling used cars. Future prediction is that there would be a huge market of used cars in the US, Canada, Mexico.

Prediction of purchase of used cars in future



## Project Architecture



We are going to use big data technologies such as Jupyter Notebook to include datasets of American cars and Canadian cars in the HDFS to do analysis of data and predict the value of used cars in the market. There is a plethora of factors that play a major role in deciding the resale value of a used car. In this project we will predict which factors affect the cost of cars, what is the most profitable car and queries as such.

## DATA CLEANING:

- Import CSV files using Pandas.

```
Used cars in US and Canada Analysis

In [1]: import numpy as np
import pandas as pd

Import US Dataset

In [6]: df=pd.read_csv(r"C:\Users\mahat\OneDrive\Desktop\big data project\Data\us-dealers-used.csv", low_memory=False)

In [7]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7184384 entries, 0 to 7184383
Data columns (total 21 columns):
#   Column      Dtype
---  -
0   id           object
1   vin          object
2   price       float64
3   miles       float64
4   stock_no    object
..
```

Count and drop all null values.

Count all null values in each row

```
In [8]: num = df.isnull().sum()
print(num)
```

```
id          0
vin         0
price      656779
miles      67290
stock_no   172288
year       160
make        0
model      8040
trim       15169
body_type  27978
vehicle_type 38366
drivetrain 15921
transmission 13600
fuel_type  45956
engine_size 103266
engine_block 107110
seller_name 11111
street     31145
city      11404
state     11414
zip       11647
```

Replace the null values in prices with its mean value. And drop null values in other columns.

## Data Cleaning

```
In [38]: # Copy "price" column into an other variable
price_column = df["price"].copy()

# Drop all rows with missing values in all columns except "price"
df = df.dropna(subset=df.columns.difference(["price"]))
df["price"] = price_column

# Calculate the mean value of non-missing values in "price"
price_mean = df["price"].mean(skipna=True)

# Replace missing values in "price" with the mean of non-missing values
df["price"] = df["price"].fillna(value=price_mean)

# Count the resulting DataFrame null values
num = df.isnull().sum()

print(num)
```

```
id          0
vin         0
price       0
miles       0
stock_no    0
year        0
make        0
model       0
trim        0
body_type   0
vehicle_type 0
drivetrain  0
transmission 0
fuel_type   0
engine_size  0
engine_block 0
seller_name  0
street      0
city        0
state       0
zip         0
dtype: int64
```

After data cleaning is completed for both data sets download the cleaned CSV file

## Download cleaned datasets in CSV format

```
In [34]: df.to_csv("us_cars.csv")
df1.to_csv("canada_cars.csv")
```

## Data Loading into Spark

```
mahathi@mahathl-VirtualBox: /usr/share/spark
```

```
>>> df = spark.read.csv("file:///home/mahathl/Downloads/us_cars.csv",inferSchema='true',header='true')
>>> df.count()
6750068
>>> df.show(5)
23/04/17 23:08:45 WARN CSVHeaderChecker: CSV header does not conform to the schema.
Header: , id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, eng
ine_block, seller_name, street, city, state, zip
Schema: _c0, id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, eng
ine_block, seller_name, street, city, state, zip
Expected: _c0 but found:
CSV file: File:///home/mahathl/Downloads/us_cars.csv
```

_c0	id	vin	price	miles	stock_no	year	make	model	trim	body_type	vehicle_type	drivetrain
[Transmission]		fuel_type	[engine_size]	[engine_block]		seller_name		street		city	[state]	zip
0 38b2f52e-8f5d 1GCWGFCF3F1284719 Automatic E85   Unleaded 4.8 20998.0 115879.0 W1T503168C 2015.0 Chevrolet Express Cargo Work Van Cargo Van Truck RWD						Vnissan ellcott city 8569 Baltimore Na... Ellcott City MD 21043						
1 97ba4955-ccf0 WBVY7Z8CSJVB87514 Automatic Electric   Premium 0.6 27921.0 7339.0 P33243 2018.0 BMW i3 S Hatchback Car RWD						I hendrick honda po... 5381 N Federal Hl... Pompano Beach FL 33064						
2 be1da9fd-0f34 ML32F4F32JHF10325 Automatic Unleaded 1.2 11055.0 39798.0 WM2091A 2018.0 Mitsubishi Mirage G4 SE Sedan Car FWD						I  russ darrow toyota 2708 West Washing... West Bend WI 53895						
3 8a77e45-6cb6 1GCPTEE15K1291189 Automatic Diesel 2.8 52997.0 28568.0 9U2Y425A 2019.0 Chevrolet Colorado ZR2 Pickup Truck 4WD						I  young kia 308 North Main St... Layton UT 84041						
4 cde691c3-91dd 1G2AL18F0R087312093 Automatic Unleaded 2.2 27889.284844765512 188485.0 T36625A 2008.0 Pontiac G5 Base Coupe Car FWD						pappas toyota 10011 Spencer Rd Saint Peters MO 63376						

only showing top 5 rows

[illegible]

## Merge both datasets, using UNION

```

>>> df = df.union(df1)
>>> df.count()
7011263
>>> df.show(5)
23/04/17 23:18:24 WARN CSVHeaderChecker: CSV header does not conform to the schema.
Header: | id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Schema: _c0, id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Expected: _c0 but found:
CSV file: File:///home/mahathi/Downloads/us_cars.csv

```

	_c0	id	vin	price	miles	stock_no	year	make	model	trim	body_type	vehicle_type	drivetrain	transmission	fuel_type	engine_size	engine_block	seller_name	street	city	state	zip		
0	38b2f52e-8f5d	1	1CGWGF3F31284719	20998.0	115879.0	W1T503168C	2015.0	Chevrolet	Express Cargo	Work Van	Cargo Van	Truck	RWD	Automatic	E85	Unleaded	4.8	V1nssan ellicott	city	8569 Baltimore Na...	Ellicott City	MD	21043	
1	97ba4955-ccf0	1	MBW7Z8C59JV887514	27921.0	7339.0	P33243	2018.0	BMW	i3	s	Matchback	Car	RWD	Automatic	Electric	Premiu...	0.6	I hendrick honda po...	5381 N Federal Hl...	Pompano Beach	FL	33064		
2	beida9fd-0f34	1	JML32F4FJ23JH10325	11055.0	39798.0	WM2091A	2018.0	Mitsubishi	Mirage G4	SE	Sedan	Car	FWD	Automatic	Unleaded	1.2	I  russ darrow toyota	2700 West Washing...	West Bend	WI	53095			
3	84327e45-6cb0	1	1GCPTEE15K1291189	25997.0	28568.0	9U2Y425A	2019.0	Chevrolet	Colorado	ZR2	Pickup	Truck	4WD	Automatic	Diesel	2.8	I  young ksl308 North Main St...	Layton	UT	84041				
4	1c1d10e1-91dd	1	1G2AL18F08731293	28489.44	765512	188485.0	2008.0	Pontiac	G5	Base	Coupe	Car	FWD	Automatic	Unleaded	2.2	I  pappas toyota	10011 Spencer Rd	Saint Peters	MO	63376			

only showing top 5 rows

Clean the data by using only relevant columns

```
mahathi@mahathi-VirtualBox: /usr/share/spark
>>> apt_column = ['price', 'miles', 'year', 'make', 'model', 'trim', 'body_type', 'vehicle_type', 'seller_name', 'city', 'state']
>>> df = df.select(*apt_column)
>>> df.printSchema()
root
 |-- price: double (nullable = true)
 |-- miles: double (nullable = true)
 |-- year: double (nullable = true)
 |-- make: string (nullable = true)
 |-- model: string (nullable = true)
 |-- trim: string (nullable = true)
 |-- body_type: string (nullable = true)
 |-- vehicle_type: string (nullable = true)
 |-- seller_name: string (nullable = true)
 |-- city: string (nullable = true)
 |-- state: string (nullable = true)
```

## Business Questions

### QUERY 1:

- Top 5 most expensive cars and their Mileage.
- Top 5 cars with the maximum Mileage and their respective prices.

```
mahathi@mahathi-VirtualBox: /usr/share/spark
>>> df.groupBy("seller_name", "vehicle_type", "make", "state", "miles").mean("price").sort("avg(price)", ascending=False).show(5)
+-----+-----+-----+-----+-----+-----+
| seller_name | vehicle_type | make | state | miles | avg(price) |
+-----+-----+-----+-----+-----+
| mclaren charlotte | Car | Porsche | NC | 5984.0 | 1499996.0 |
| ferrari of newpor... | Car | Ferrari | CA | 859.0 | 1495000.0 |
| karbuds | Car | Ferrari | FL | 697.0 | 1479900.0 |
| cauley ferrari of... | Car | Ferrari | MI | 697.0 | 1479900.0 |
| cauley ferrari | Car | Ferrari | MI | 697.0 | 1479900.0 |
+-----+-----+-----+-----+-----+
only showing top 5 rows

>>> df.groupBy("seller_name", "vehicle_type", "make", "state", "price").mean("miles").sort("avg(miles)", ascending=False).show(5)
+-----+-----+-----+-----+-----+-----+
| seller_name | vehicle_type | make | state | price | avg(miles) |
+-----+-----+-----+-----+-----+
| 3 sons auto sales | Truck | Ford | IL | 12990.0 | 3000000.0 |
| 3 sons auto sales | Truck | Acura | IL | 7990.0 | 3000000.0 |
| cleveland motor c... | Truck | Chevrolet | OH | 27889.284844765512 | 2975291.0 |
| lexus of clearwater | Truck | Ford | FL | 45989.0 | 2575500.0 |
| lexus of tampa bay | Truck | Ford | FL | 45989.0 | 2575500.0 |
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

By observing the results, we can conclude that the most expensive cars drive very less miles, and their market share is dominated by Ferrari cars. If the expectation of the vehicle is to give more miles and costs less than Trucks are preferable and Ford is making most of them.

## QUERY 2:

- Market Share of all companies

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> from pyspark.sql.functions import count, sum, col
>>> total_cars = df.count()
>>> car_counts = df.groupBy("make").agg(count("*").alias("car_count"))
>>> market_share = car_counts.withColumn("market_share", col("car_count") / total_cars * 100).orderBy(col("market_share").desc())
>>> market_share.show()
+-----+-----+-----+
| make | car_count | market_share |
+-----+-----+-----+
| Ford | 917695 | 13.088868581880325 |
| Chevrolet | 791692 | 11.291717341083912 |
| Toyota | 639819 | 9.125588356905167 |
| Honda | 514314 | 7.335539973325776 |
| Nissan | 448512 | 6.397021478155933 |
| Jeep | 427022 | 6.0905146476462235 |
| Hyundai | 278517 | 3.9724226576581136 |
| BMW | 266249 | 3.7974470505528033 |
| GMC | 254194 | 3.6255094124981477 |
| Kia | 231725 | 3.3050393345678235 |
| Mercedes-Benz | 219127 | 3.1253570148488223 |
| RAM | 211126 | 3.011240628115077 |
| Dodge | 203757 | 2.9061383091748234 |
| Subaru | 186700 | 2.6628583181090195 |
| Lexus | 177118 | 2.5261924991260494 |
| Volkswagen | 163148 | 2.3269416651464936 |
| Audi | 125041 | 1.7834304603892337 |
| Mazda | 117525 | 1.6762315149210634 |
| Cadillac | 100415 | 1.43219559728397 |
| Acura | 98537 | 1.4054101236824235 |
+-----+-----+-----+
only showing top 20 rows
```

We can observe that Ford, Chevrolet, and Toyota take over 30% of the market share.

## QUERY 3

- Cars with the maximum market share over the time

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> from pyspark.sql.functions import rank
>>> from pyspark.sql.window import Window
>>> market_share_df = df.groupBy("make", "year").agg(sum("price").alias("total_sales"))
>>> total_sales_df = market_share_df.groupBy("year").agg(sum("total_sales").alias("total_sales_year"))
>>> market_share_df = market_share_df.join(total_sales_df, "year")
>>> market_share_df = market_share_df.withColumn("market_share", market_share_df["total_sales"] / market_share_df["total_sales_year"])
>>> market_share_df = market_share_df.drop("total_sales_year")
>>> window = Window.partitionBy("year").orderBy(col("market_share").desc())
>>> highest_market_share_df = market_share_df.select("year", rank().over(window).alias("rank")).filter("rank = 1").orderBy(col("year").desc())
>>> highest_market_share_df.show()
+-----+-----+-----+-----+
| year | make | total_sales | market_share | rank |
+-----+-----+-----+-----+
| 2022 | Acura | 741339.0 | 0.5762690388869801 | 1 |
| 2021 | BMW | 6.0415764E7 | 0.18083605850610515 | 1 |
| 2020 | Ford | 1.28449534E8 | 0.11803152348112898 | 1 |
| 2019 | Ford | 1.75805011E8 | 0.12680506700451305 | 1 |
| 2018 | Ford | 1.52509477E8 | 0.10917571936817914 | 1 |
| 2017 | Ford | 1.16771678E8 | 0.09709864842935285 | 1 |
| 2016 | Ford | 5.9067879E7 | 0.10168036494560344 | 1 |
| 2015 | Ford | 3.8696919E7 | 0.1161150294535269 | 1 |
| 2014 | Ford | 2.9019577E7 | 0.13448475682502942 | 1 |
| 2013 | Ford | 2.2517639E7 | 0.15397808781053388 | 1 |
| 2012 | Ford | 9583555.0 | 0.10927574921132005 | 1 |
| 2011 | Ford | 8059224.0 | 0.13604589043527415 | 1 |
| 2010 | Ford | 5789767.0 | 0.1468610075334358 | 1 |
| 2009 | Toyota | 2782296.0 | 0.10447235284257729 | 1 |
| 2008 | Ford | 3085392.0 | 0.13136248506074627 | 1 |
| 2007 | Ford | 1893850.0 | 0.11788251299113409 | 1 |
| 2006 | Chevrolet | 1282363.0 | 0.12225757451096274 | 1 |
| 2005 | Ford | 571554.0 | 0.1011453026029315 | 1 |
| 2004 | Ford | 691316.0 | 0.1403028954332292 | 1 |
| 2003 | Ferrari | 855314.0 | 0.2430659282448614 | 1 |
+-----+-----+-----+-----+
only showing top 20 rows
```

Market share of each car over the time

The car industry is highly competitive, with many different makes of cars vying for market share.



Ford appears to be a consistently popular make, with sales in the top rank for many years. This could be due to factors such as brand loyalty, affordability, or a wide range of available models.

Other makes such as Acura, BMW, Chevrolet, Ferrari, and Toyota have also had strong sales in certain years.

The data covers a period of 20 years, from 2003 to 2022. Looking at trends over time could provide insights into changes in consumer preferences, economic conditions, or other factors that affect car sales.

The data can also be used to compare the performance of different car makes and identify patterns in sales. For example, some makes may have experienced a decline in sales over time, while others may have seen a steady increase.

### QUERY 4

- Correlation between “price” and other factors. Which factors affect the price of cars.

```
mahathi@mahathi-VirtualBox: /usr/share/spark
>>> from pyspark.sql.functions import corr
>>> df.select(corr("price", "miles"),
...          corr("price", "year"),
...          corr("price", "make"),
...          corr("price", "model"),
...          corr("price", "trim"),
...          corr("price", "body_type"),
...          corr("price", "vehicle_type"),
...          corr("price", "seller_name"),
...          corr("price", "city"),
...          corr("price", "state")).show()
23/04/18 02:22:40 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.
+-----+
| corr(price, miles)| corr(price, year)| corr(price, make)| corr(price, model)| corr(price, trim)| corr(price, body_type)| corr(price, vehicle_type)| corr(price, seller_name)| corr(price, city)|
| corr(price, state)|
+-----+
|-0.39453844575817854| 0.4204254553788824| null| 0.5259314566474526| 0.3964831637921187| null| null| null| null|
| null|
+-----+
```

This table shows the correlation coefficients between the price column and each of the other columns in the dataset. There is a moderate negative correlation between price and miles (-0.39) and a moderate positive correlation between price and year (0.42). The correlation coefficients indicate that as the mileage of a used car increases, its resale value tends to decrease. There is also a moderate positive correlation between price and model (0.53) and trim (0.40). However, there is no correlation between price and make, body\_type, vehicle\_type, seller\_name, city, or state it is to the contrary belief to see that the make of the model plays no role in the resale value of the car.

## QUERY 5

- Average price and mileage of cars in descending order.

```

mahathi@mahathi-VirtualBox: /usr/share/spark
>>> from pyspark.sql.functions import avg
>>> df.groupBy('body_type', 'vehicle_type').agg(avg('price').alias('avg_price'), avg('miles').alias('avg_mileage')).orderBy(col("avg_price").desc()).show()
+-----+-----+-----+-----+
| body_type|vehicle_type| avg_price| avg_mileage|
+-----+-----+-----+-----+
| Targa|Car|90393.29837207461|16112.233151803948|
| Convertible|Car|42803.842576089446|46131.58430221585|
| Roadster|Car|40319.49759893186|42645.9806324239|
| Chassis Cab|Truck|38839.8586236172|80770.40039361469|
| Pickup|Truck|37762.54774964416|60611.45363272988|
| Coupe|Car|37144.91974477872|45036.341911445015|
| Cutaway|Truck|34179.59007675749|62138.87579393084|
| Combi|Truck|32862.31548023485|37447.54384133612|
| SUV|Truck|28664.655819943084|52916.044624374386|
| Sedan|Truck|27889.284844765512|112444.0|
| Cargo Van|Truck|27884.268035456862|55824.589241850204|
| Passenger Van|Truck|27218.77891922609|65716.39552919708|
| Wagon|Car|23935.940556910828|57583.2343285676|
| Crossover|Truck|23907.845813081614|44822.950457599836|
| Minivan|Truck|22000.65040204023|78364.44849522546|
| Sedan|Car|21486.6736575073|56733.799486144|
| Crossover|Car|21441.95918465987|27674.301040163995|
| Hatchback|Car|19364.1764448648|53997.2282544099|
| Car Van|Truck|18621.382981840703|89043.42561205273|
| SUV|Car|17808.70500012859|63104.34043701106|
+-----+-----+-----+-----+
only showing top 20 rows

```

The output shows the average price and mileage for each body type and vehicle type. The most expensive car is "Targa" is \$90,393.30 and the average mileage is 16,112.23. The table provides insights into how prices and mileage vary by body type and vehicle type, which can be useful for understanding market trends and making informed purchasing decisions.

## PREDICTION:

```

mahathi@mahathi-VirtualBox: /usr/share/spark
>>> from pyspark.ml.regression import RandomForestRegressor
>>> from pyspark.ml.feature import VectorAssembler
>>> from pyspark.ml.linalg import Vectors
>>> from pyspark.ml.evaluation import RegressionEvaluator
>>> assembler = VectorAssembler(inputCols=["miles", "year"], outputCol="features")
>>> data = assembler.transform(df)
>>> (trainingData, testData) = data.randomSplit([0.7, 0.3])
>>> rf = RandomForestRegressor(labelCol="price", featuresCol="features", numTrees=10, maxDepth=5, seed=42)
>>> model = rf.fit(trainingData)
23/04/20 17:37:56 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:37:56 WARN BlockManager: Persisting block rdd_140_2 to disk instead.
23/04/20 17:38:01 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 3.4 MiB so far)
23/04/20 17:38:01 WARN BlockManager: Persisting block rdd_140_3 to disk instead.
23/04/20 17:38:09 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:09 WARN BlockManager: Persisting block rdd_140_4 to disk instead.
23/04/20 17:38:16 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:16 WARN BlockManager: Persisting block rdd_140_5 to disk instead.
23/04/20 17:38:23 WARN MemoryStore: Not enough space to cache rdd_140_6 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:23 WARN BlockManager: Persisting block rdd_140_6 to disk instead.
23/04/20 17:38:30 WARN MemoryStore: Not enough space to cache rdd_140_7 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:30 WARN BlockManager: Persisting block rdd_140_7 to disk instead.
23/04/20 17:38:38 WARN MemoryStore: Not enough space to cache rdd_140_8 in memory! (computed 5.2 MiB so far)
23/04/20 17:38:38 WARN BlockManager: Persisting block rdd_140_8 to disk instead.
23/04/20 17:38:45 WARN MemoryStore: Not enough space to cache rdd_140_9 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:45 WARN BlockManager: Persisting block rdd_140_9 to disk instead.
23/04/20 17:38:49 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:51 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:52 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:54 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:55 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:58 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:00 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:01 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:02 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:04 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:08 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:09 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:10 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:12 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:13 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:18 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:19 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:22 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:23 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:25 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
>>> predictions = model.transform(testData)
>>> evaluator = RegressionEvaluator(labelCol="price", predictionCol="prediction", metricName="rmse")
>>> rmse = evaluator.evaluate(predictions)
>>> new_data = [(10000, 2025, "Ford", "Mustang", "GT", "Coupe", "New", "Dealer", "New York", "NY"),
... (50000, 2038, "Toyota", "Camry", "SE", "Sedan", "Used", "Private Seller", "Los Angeles", "CA")]
>>> new_data_df = spark.createDataFrame(new_data, ["miles", "year", "make", "model", "trin", "body_type", "vehicle_type", "seller_name", "city", "state"])
>>> assembler = VectorAssembler(inputCols=["miles", "year"], outputCol="features")
>>> new_data_features = assembler.transform(new_data_df)
>>> predictions = model.transform(new_data_features)
>>> predictions.select("make", "model", "year", "miles", "prediction").show()

```

Predicted prices of cars in future:

make	model	year	miles	prediction
Ford	Mustang	2025	10000	41127.86025576979
Toyota	Camry	2038	50000	36882.95919682251

### CONCLUSION:

Based on this study, we can draw several conclusions regarding the efficiency, market share, and pricing of different types of vehicles. Firstly, it appears that the most expensive cars tend to offer lower mileage, while trucks are both the most efficient and least expensive option, while also providing maximum mileage. Ford, Chevrolet, and Toyota are the dominant players in the market, with Honda and Nissan following closely behind in terms of sales numbers and potential for future growth. It is worth noting that Ford has consistently ranked as the top truck maker for the past two decades.

Furthermore, our correlation study has revealed that mileage has a negative impact on the price of used cars, while the year, model, and trim of the vehicle have a positive effect on the price. This suggests that dealers should prioritize acquiring recently bought and lightly used cars in order to maximize their profits.

Finally, our prediction model has a root mean square error of 4574.7. While this is not ideal, we believe that with additional data and further training, we can reduce the RSME value and achieve more accurate predictions.



---

# USED CARS

BIG DATA FINAL PROJECT



## FRAMEWORK USED



Jupyter Notebook



Hadoop

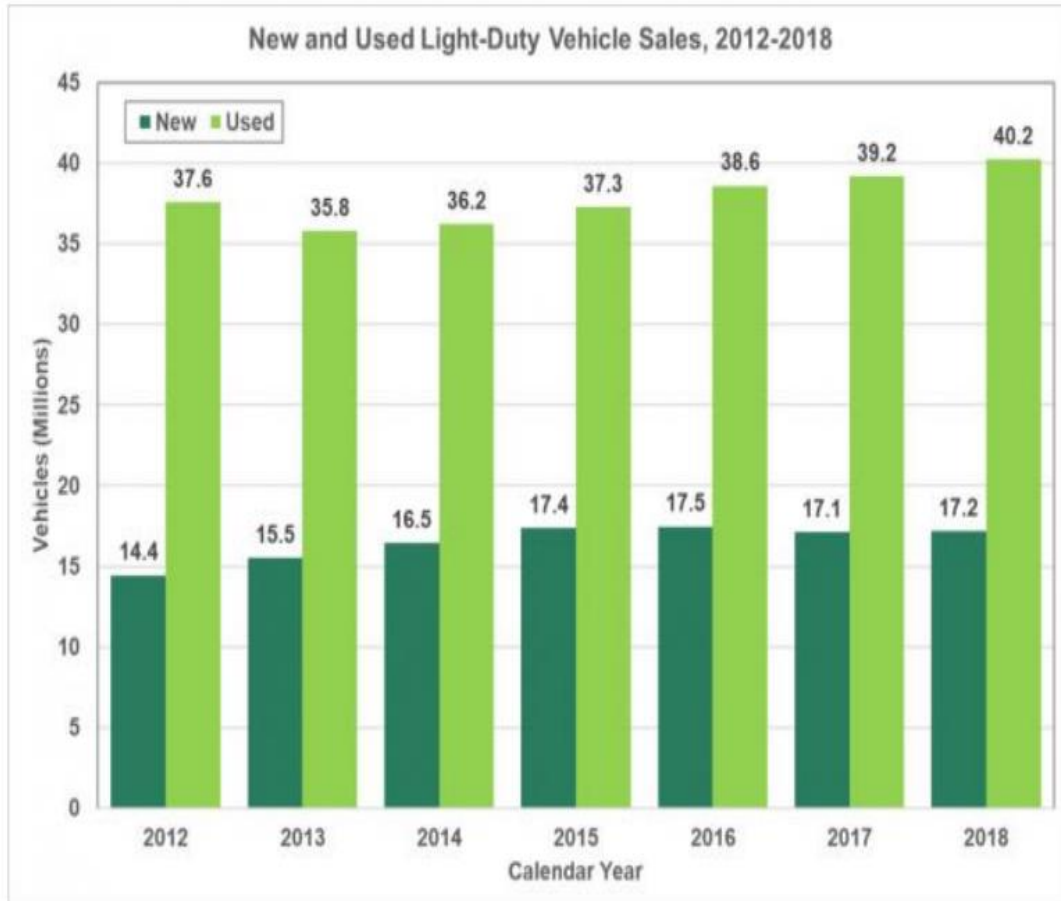


Spark

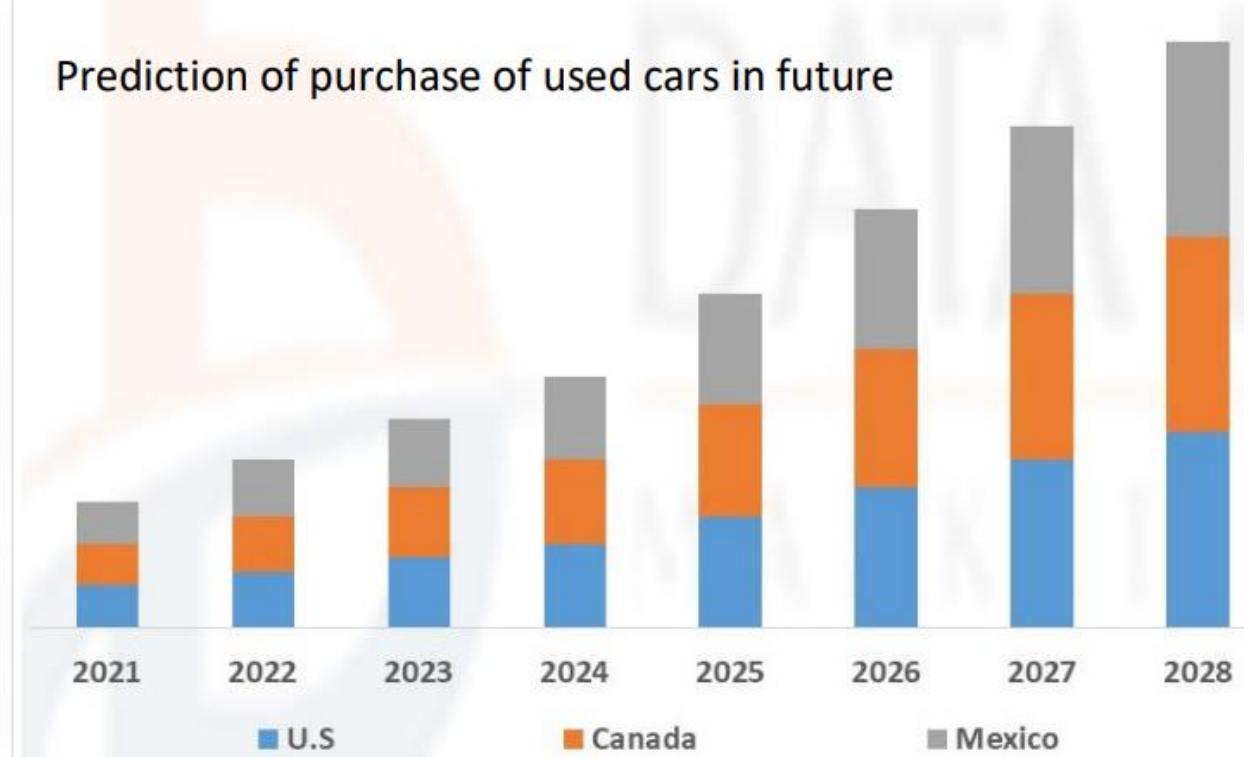
## OBJECTIVE

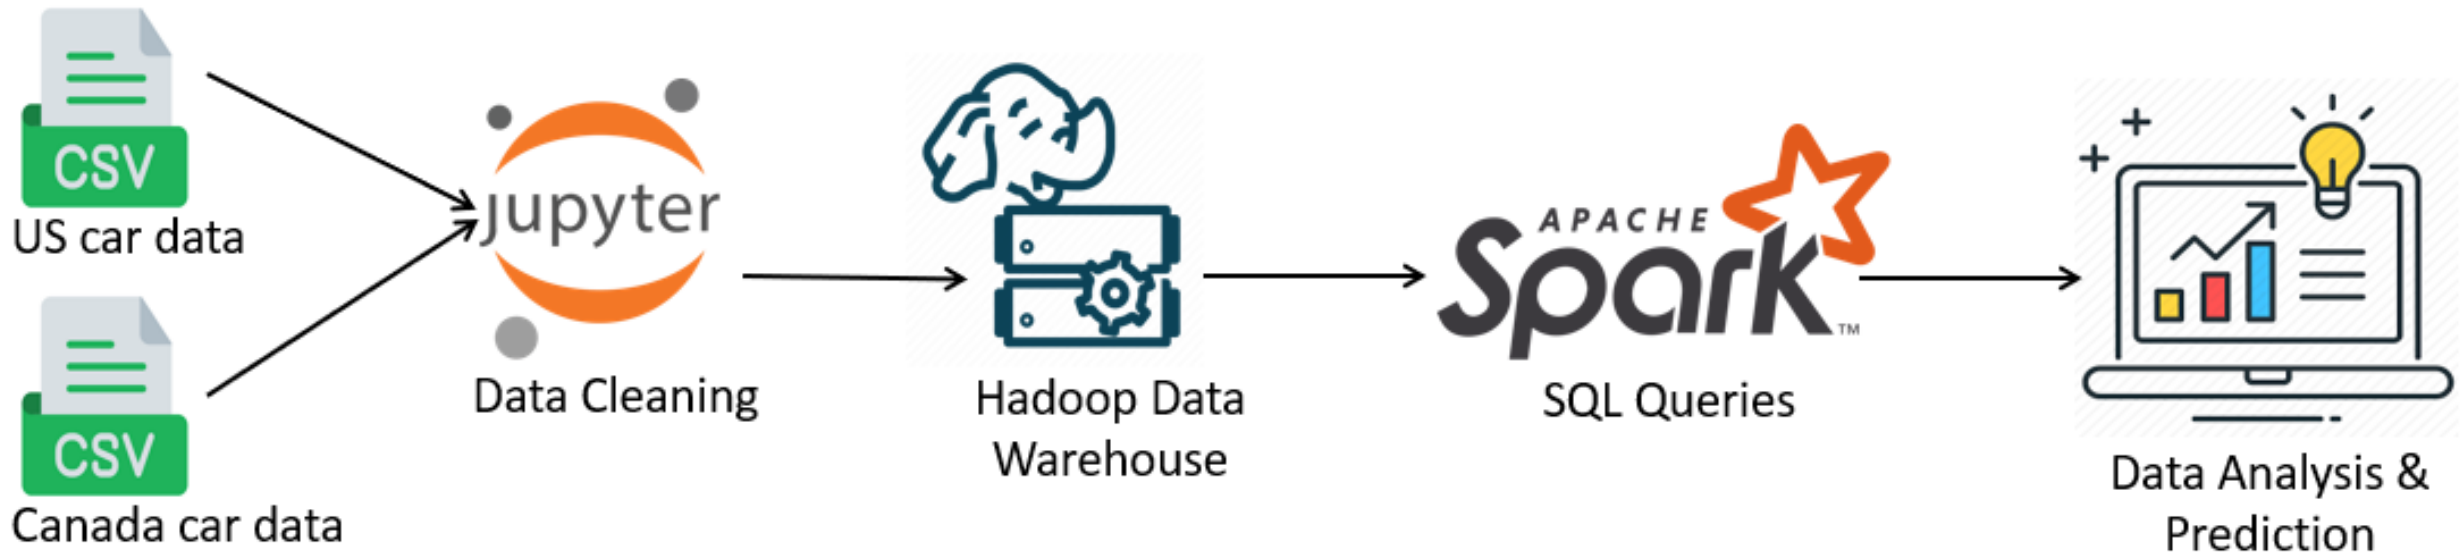
In this project, we are leveraging these technology components to:

- Perform analysis on used car market data with multi-million records
  - By loading the unstructured data to data warehouse
  - Utilize spark analytics engine to interface with HDFS and process data
- Train ML models and predict car prices using generated model
- Provide Business recommendations to the sellers to make their sale profitable



Prediction of purchase of used cars in future







# DATASET SUMMARY

Data Field	Data Description
id	This is a GUID and unique in the feed
vin	17 char long VIN of the car
price	The car price as listed on the website
miles	The car miles/odometer as listed on the website
stock_no	Stock number of the car listed on the website
year	Model Year of the car (VIN Decoded)
make	Make of the car (VIN Decoded)
model	Model of the car (VIN Decoded)
trim	Trim of the car (VIN Decoded)
vehicle_type	Vehicle type of the car (VIN Decoded)
body_type	Body type of the car (VIN Decoded)
drivetrain	Drivetrain of the car (VIN Decoded)
fuel_type	Fuel type of the car (VIN Decoded)
engine_block	Engine block of the car (VIN Decoded)
engine_size	Engine size of the car (VIN Decoded)
transmission	Transmission of the car (VIN Decoded)
seller_name	Dealer Name
city	Dealer Location
state	Dealer Location
zip	Dealer Location

There are two data sets in source zipped file:

- US used car dataset –
  - 7,104,304 records
  - 1.28 GB size
- Canada used car dataset –
  - 393,603 records
  - 70 MB size

Link to dataset zipped file:

<https://www.kaggle.com/datasets/rupesthaurandal/marketcheck-automotive-data-us-canada>

# DATA CLEANING

Count all null values in each row

```
In [8]: num = df.isnull().sum()  
print(num)
```

```
id          0  
vin         0  
price      656779  
miles      67290  
stock_no   172288  
year       160  
make       0  
model     8040  
trim     15169  
body_type 27978  
vehicle_type 38366  
drivetrain 15921  
transmission 13600  
fuel_type  45956  
engine_size 103266  
engine_block 107110  
seller_name 11111  
street     31145  
city      11404  
state     11414  
zip       11647
```

```
id          0  
vin         0  
price      0  
miles      0  
stock_no   0  
year       0  
make       0  
model      0  
trim       0  
body_type  0  
vehicle_type 0  
drivetrain 0  
transmission 0  
fuel_type  0  
engine_size 0  
engine_block 0  
seller_name 0  
street     0  
city       0  
state      0  
zip        0  
dtype: int64
```

# REPLACE NULL VALUES IN PRICE WITH MEAN VALUE

```
In [14]: ▶ n=len(df)
          print(n)
```

```
7104304
```

## Data Cleaning

```
In [38]: ▶ # Copy "price" column into an other variable
          price_column = df["price"].copy()

          # Drop all rows with missing values in all columns except "price"
          df = df.dropna(subset=df.columns.difference(["price"]))
          df["price"] = price_column

          # Calculate the mean value of non-missing values in "price"
          price_mean = df["price"].mean(skipna=True)

          # Replace missing values in "price" with the mean of non-missing values
          df["price"] = df["price"].fillna(value=price_mean)

          # Count the resulting DataFrame null values
          num = df.isnull().sum()

          print(num)
```

# SPARK - LOADING DATA INTO DATAFRAMES

```
mahathi@mahathi-VirtualBox: /usr/share/spark
```

```
>>> df = spark.read.csv("file:///home/mahathi/Downloads/us_cars.csv",inferSchema='true',header='true')
>>> df.count()
6750668
>>> df.show(5)
23/04/17 23:08:45 WARN CSVHeaderChecker: CSV header does not conform to the schema.
Header: , id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Schema: _c0, id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Expected: _c0 but found:
CSV file: file:///home/mahathi/Downloads/us_cars.csv
```

_c0	id	vin	price	miles	stock_no	year	make	model	trim	body_type	vehicle_type	drivetrain
[transmission]			fuel_type engine_size engine_block		seller_name		street		city state  zip			
0 38b2f52e-8f5d 1GCWGF3F1284719			20998.0	115879.0	W1T503168C	2015.0	Chevrolet	Express Cargo	Work Van	Cargo Van	Truck	RWD
Automatic	E85 / Unleaded	4.8	V	nissan ellicott city	8569 Baltimore Na...	Ellicott City	MD	21043				
1 97ba4955-ccf0 WBVY7Z8C59JV87514			27921.0	7339.0	P33243	2018.0	BMW	i3	s	Hatchback	Car	RWD
Automatic	Electric / Prentiu...	0.6	I hendrick honda po...	5381 N Federal Hl...	Pompano Beach	FL	33064					
2 be1da9fd-0f34 ML32FAF32JWF10325			11055.0	39798.0	WM2091A	2018.0	Mitsubishi	Mirage G4	SE	Sedan	Car	FWD
Automatic	Unleaded	1.2	I  russ darrow toyota	2700 West Washing...	West Bend	WI	53095					
3 84327e45-6cb6 1GCPTEE15K1291189			52997.0	28568.0	9U2Y425A	2019.0	Chevrolet	Colorado	ZR2	Pickup	Truck	4WD
Automatic	Diesel	2.8	I  young kia	308 North Main St...	Layton	UT	84041					
4 cde691c3-91dd 1G2AL18F087312093			27889.284844765512	188485.0	T36625A	2008.0	Pontiac	G5	Base	Coupe	Car	FWD
Automatic	Unleaded	2.2	I  pappas toyota	10011 Spencer Rd	Saint Peters	MO	63376					

```
only showing top 5 rows
```

```
mahathi@mahathi-VirtualBox: /usr/share/spark
```

```
>>> df1 = spark.read.load("file:///home/mahathi/Downloads/canada_cars.csv",format='csv',inferSchema='true',header='true')
>>> df1.show(5)
23/04/17 23:12:48 WARN CSVHeaderChecker: CSV header does not conform to the schema.
Header: , id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Schema: _c0, id, vin, price, miles, stock_no, year, make, model, trin, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Expected: _c0 but found:
CSV file: file:///home/mahathi/Downloads/canada_cars.csv
```

_c0	id	vin	price	miles	stock_no	year	make	model	trim	body_type	vehicle_type	drivetrain	transmission		fuel
	_type engine_size engine_block		seller_name							city state	zip				
0 b39ea795-eca9	19UNC1B01HY800062	179999.0	9966.0	V-P4139	2017.0	Acura	NSX Base	Coupe	Car	4WD	Automatic Electric / Prem				
iu...	3.5	V	edmundston honda	475 Rue Victoria	Edmundston	NB E3V 2K7									
1 026cb5b1-6e3e	19UNC1B02HY800023	179995.0	5988.0	PPAP70374	2017.0	Acura	NSX Base	Coupe	Car	4WD	Automatic Electric / Prem				
iu...	3.5	V garage daniel les...	2795 Route-du-prs...	Notre-dame-des-pins	QC G0M 1K0										
2 5cd5d5b2-5cc2	19UNC1B02HY800071	168528.0	24242.0	B21085	2017.0	Acura	NSX Base	Coupe	Car	4WD	Automatic Electric / Prem				
iu...	3.5	V	lougheed acura 1388 Lougheed Hlg...	Coquitlam	BC V3K 6S4										
3 b32473ed-5922	19UNC1B02LY800001	220000.0	6637.0	AP5333	2020.0	Acura	NSX Base	Coupe	Car	4WD	Automatic Electric / Prem				
iu...	3.5	V	drive autogroup 1305 Parkway Suit...	Pickering	ON L1V 3P2										
4 ac40c9fc-0676	19UNC1B02LY800001	220000.0	6637.0	AP5333	2020.0	Acura	NSX Base	Coupe	Car	4WD	Automatic Electric / Prem				
iu...	3.5	V	acura pickering	575 Kingston Road	Pickering	ON L1V 3N7									

```
only showing top 5 rows
```

# MERGE DATASETS

```
>>> df = df.union(df1)
>>> df.count()
7011263
>>> df.show(5)
23/04/17 23:18:24 WARN CSVHeaderChecker: CSV header does not conform to the schema.
Header: , id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Schema: _c0, id, vin, price, miles, stock_no, year, make, model, trim, body_type, vehicle_type, drivetrain, transmission, fuel_type, engine_size, engine_block, seller_name, street, city, state, zip
Expected: _c0 but found:
CSV file: file:///home/mahathi/Downloads/us_cars.csv
```

_c0	id	vin	price	miles	stock_no	year	make	model	trim	body_type	vehicle_type	drivetrain
transmission		fuel_type	engine_size	engine_block	seller_name		street		city	state	zip	
0 38b2f52e-8f5d 1GCGWFCF3F1284719			20998.0	115879.0	W1T503168C	2015.0	Chevrolet	Express Cargo	Work Van	Cargo Van	Truck	RWD
Automatic	E85 / Unleaded	4.8		V nissan ellicott city	8569 Baltimore Na...	Ellicott City	MD	21043				
1 97ba4955-ccf0 WBYYZ8CS9JVB87514			27921.0	7339.0	P33243	2018.0	BMW	i3	s	Hatchback	Car	RWD
Automatic	Electric / Premium	0.6		I hendrick honda po...	5381 N Federal Hl...	Pompano Beach	FL	33064				
2 be1da9fd-0f34 ML32F4FJ2JHF10325			11055.0	39798.0	WM2091A	2018.0	Mitsubishi	Mirage G4	SE	Sedan	Car	FWD
Automatic	Unleaded	1.2		I  russ darrow toyota	2700 West Washing...	West Bend	WI	53095				
3 84327e45-6cb6 1GCPTEE15K1291189			52997.0	28568.0	9U2Y425A	2019.0	Chevrolet	Colorado	ZR2	Pickup	Truck	4WD
Automatic	Diesel	2.8		I  young kia	308 North Main St...	Layton	UT	84041				
4 cde691c3-91dd 1G2AL18F087312093			27889.284844765512	188485.0	T36625A	2008.0	Pontiac	G5	Base	Coupe	Car	FWD
Automatic	Unleaded	2.2		I  pappas toyota	10011 Spencer Rd	Saint Peters	MO	63376				

only showing top 5 rows

# RELEVANT COLUMNS

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> apt_column = ['price', 'miles', 'year', 'make', 'model', 'trim', 'body_type', 'vehicle_type', 'seller_name', 'city', 'state']
>>> df = df.select(*apt_column)
>>> df.printSchema()
root
 |-- price: double (nullable = true)
 |-- miles: double (nullable = true)
 |-- year: double (nullable = true)
 |-- make: string (nullable = true)
 |-- model: string (nullable = true)
 |-- trim: string (nullable = true)
 |-- body_type: string (nullable = true)
 |-- vehicle_type: string (nullable = true)
 |-- seller_name: string (nullable = true)
 |-- city: string (nullable = true)
 |-- state: string (nullable = true)
```

# PRICES VS MILEAGE

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> df.groupBy("seller_name","vehicle_type","make","state","miles").mean("price").sort("avg(price)", ascending=False).show(5)
+-----+-----+-----+-----+-----+-----+
| seller_name | vehicle_type | make | state | miles | avg(price) |
+-----+-----+-----+-----+-----+
| mclaren charlotte | Car | Porsche | NC | 5984.0 | 1499996.0 |
| ferrari of newpor... | Car | Ferrari | CA | 859.0 | 1495000.0 |
| karbuds | Car | Ferrari | FL | 697.0 | 1479900.0 |
| cauley ferrari of... | Car | Ferrari | MI | 697.0 | 1479900.0 |
| cauley ferrari | Car | Ferrari | MI | 697.0 | 1479900.0 |
+-----+-----+-----+-----+-----+
only showing top 5 rows

>>> df.groupBy("seller_name","vehicle_type","make","state","price").mean("miles").sort("avg(miles)", ascending=False).show(5)
+-----+-----+-----+-----+-----+-----+
| seller_name | vehicle_type | make | state | price | avg(miles) |
+-----+-----+-----+-----+-----+
| 3 sons auto sales | Truck | Ford | IL | 12990.0 | 3000000.0 |
| 3 sons auto sales | Truck | Acura | IL | 7990.0 | 3000000.0 |
| cleveland motor c... | Truck | Chevrolet | OH | 27889.284844765512 | 2975291.0 |
| lexus of clearwater | Truck | Ford | FL | 45989.0 | 2575500.0 |
| lexus of tampa bay | Truck | Ford | FL | 45989.0 | 2575500.0 |
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

# MARKET SHARE OF CARS

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> from pyspark.sql.functions import count, sum, col
>>> total_cars = df.count()
>>> car_counts = df.groupBy("make").agg(count("*").alias("car_count"))
>>> market_share = car_counts.withColumn("market_share", col("car_count") / total_cars * 100).orderBy(col("market_share").desc())
>>> market_share.show()
+-----+-----+-----+
|      make|car_count|market_share|
+-----+-----+-----+
|      Ford|  917695|13.088868581880325|
| Chevrolet|  791692|11.291717341083912|
|      Toyota|  639819| 9.125588356905167|
|      Honda|  514314| 7.335539973325776|
|      Nissan|  448512| 6.397021478155933|
|      Jeep|  427022|6.0905146476462235|
|     Hyundai|  278517|3.9724226576581136|
|      BMW|  266249|3.7974470505528033|
|      GMC|  254194|3.6255094124981477|
|      Kia|  231725|3.3050393345678235|
|Mercedes-Benz|  219127|3.1253570148488223|
|      RAM|  211126| 3.011240628115077|
|      Dodge|  203757|2.9061383091748234|
|      Subaru|  186700|2.6628583181090195|
|      Lexus|  177118|2.5261924991260494|
| Volkswagen|  163148|2.3269416651464936|
|      Audi|  125041|1.7834304603892337|
|      Mazda|  117525|1.6762315149210634|
|    Cadillac|  100415| 1.43219559728397|
|      Acura|   98537|1.4054101236824235|
+-----+-----+-----+
only showing top 20 rows
```



# MARKET SHARE OVER A PERIOD OF TIME

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> from pyspark.sql.functions import rank
>>> from pyspark.sql.window import Window
>>> market_share_df = df.groupBy("make", "year").agg(sum("price").alias("total_sales"))
>>> total_sales_df = market_share_df.groupBy("year").agg(sum("total_sales").alias("total_sales_year"))
>>> market_share_df = market_share_df.join(total_sales_df, "year")
>>> market_share_df = market_share_df.withColumn("market_share", market_share_df["total_sales"]/market_share_df["total_sales_year"])
>>> market_share_df = market_share_df.drop("total_sales_year")
>>> window = Window.partitionBy("year").orderBy(col("market_share").desc())
>>> highest_market_share_df = market_share_df.select("year", rank().over(window).alias("rank")).filter("rank = 1").orderBy(col("year").desc())
>>> highest_market_share_df.show()
+-----+-----+-----+-----+-----+
| year|    make| total_sales|    market_share|rank|
+-----+-----+-----+-----+-----+
|2022.0|    Acura|   741339.0| 0.5762690388869801| 1|
|2021.0|     BMW| 6.0415764E7|0.18083605850610515| 1|
|2020.0|     Ford|1.28449534E8|0.11803152348112898| 1|
|2019.0|     Ford|1.75805011E8|0.12680506700451305| 1|
|2018.0|     Ford|1.52509477E8|0.10917571936817914| 1|
|2017.0|     Ford|1.16771678E8|0.09709864842935285| 1|
|2016.0|     Ford| 5.9067879E7|0.10168036494560344| 1|
|2015.0|     Ford| 3.8696919E7| 0.1161150294535269| 1|
|2014.0|     Ford| 2.9019577E7|0.13448475682502942| 1|
|2013.0|     Ford| 2.2517639E7|0.15397808781053388| 1|
|2012.0|     Ford|   9583555.0|0.10927574921132005| 1|
|2011.0|     Ford|  8059224.0|0.13604589043527415| 1|
|2010.0|     Ford|  5789767.0| 0.1468610075334358| 1|
|2009.0|   Toyota| 2782296.0|0.10447235284257729| 1|
|2008.0|     Ford| 3085392.0|0.13136248506074627| 1|
|2007.0|     Ford| 1893850.0|0.11788251299113409| 1|
|2006.0|Chevrolet| 1282363.0|0.12225757451096274| 1|
|2005.0|     Ford|   571554.0| 0.1011453026029315| 1|
|2004.0|     Ford|   691316.0| 0.1403028954332292| 1|
|2003.0|   Ferrari|   855314.0| 0.2430659282448614| 1|
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

# CORRELATION BETWEEN PRICE AND OTHER FACTORS

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> from pyspark.sql.functions import corr
>>> df.select(corr("price", "miles"),
...          corr("price", "year"),
...          corr("price", "make"),
...          corr("price", "model"),
...          corr("price", "trim"),
...          corr("price", "body_type"),
...          corr("price", "vehicle_type"),
...          corr("price", "seller_name"),
...          corr("price", "city"),
...          corr("price", "state")).show()
23/04/18 02:22:40 WARN package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.
+-----+
| corr(price, miles)| corr(price, year)|corr(price, make)|corr(price, model)| corr(price, trim)|corr(price, body_type)|corr(price, vehicle_type)|corr(price, seller_name)|corr(price, city)|
|corr(price, state)|
+-----+
|-0.39453844575817054|0.4204254553788824|          null|0.5259314566474526|0.3964831637921107|          null|          null|          null|          null|
|          null|
+-----+
```

# RANDOM FOREST REGRESSION MODEL

```
mahathi@mahathi-VirtualBox: /usr/share/spark

>>> from pyspark.ml.regression import RandomForestRegressor
>>> from pyspark.ml.feature import VectorAssembler
>>> from pyspark.ml.linalg import Vectors
>>> from pyspark.ml.evaluation import RegressionEvaluator
>>> assembler = VectorAssembler(inputCols=["miles", "year"], outputCol="features")
>>> data = assembler.transform(df)
>>> (trainingData, testData) = data.randomSplit([0.7, 0.3])
>>> rf = RandomForestRegressor(labelCol="price", featuresCol="features", numTrees=10, maxDepth=5, seed=42)
>>> model = rf.fit(trainingData)
23/04/20 17:37:56 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:37:56 WARN BlockManager: Persisting block rdd_140_2 to disk instead.
23/04/20 17:38:01 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 3.4 MiB so far)
23/04/20 17:38:01 WARN BlockManager: Persisting block rdd_140_3 to disk instead.
23/04/20 17:38:09 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:09 WARN BlockManager: Persisting block rdd_140_4 to disk instead.
23/04/20 17:38:16 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:16 WARN BlockManager: Persisting block rdd_140_5 to disk instead.
23/04/20 17:38:23 WARN MemoryStore: Not enough space to cache rdd_140_6 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:23 WARN BlockManager: Persisting block rdd_140_6 to disk instead.
23/04/20 17:38:30 WARN MemoryStore: Not enough space to cache rdd_140_7 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:30 WARN BlockManager: Persisting block rdd_140_7 to disk instead.
23/04/20 17:38:38 WARN MemoryStore: Not enough space to cache rdd_140_8 in memory! (computed 5.2 MiB so far)
23/04/20 17:38:38 WARN BlockManager: Persisting block rdd_140_8 to disk instead.
23/04/20 17:38:45 WARN MemoryStore: Not enough space to cache rdd_140_9 in memory! (computed 17.6 MiB so far)
23/04/20 17:38:45 WARN BlockManager: Persisting block rdd_140_9 to disk instead.
23/04/20 17:38:49 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:51 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:52 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:54 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:55 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
23/04/20 17:38:58 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:00 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:01 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:02 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:04 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:08 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:09 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:10 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:12 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:13 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:18 WARN MemoryStore: Not enough space to cache rdd_140_1 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:19 WARN MemoryStore: Not enough space to cache rdd_140_2 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:22 WARN MemoryStore: Not enough space to cache rdd_140_3 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:23 WARN MemoryStore: Not enough space to cache rdd_140_4 in memory! (computed 63.2 MiB so far)
23/04/20 17:39:25 WARN MemoryStore: Not enough space to cache rdd_140_5 in memory! (computed 63.2 MiB so far)
>>> predictions = model.transform(testData)
>>> evaluator = RegressionEvaluator(labelCol="price", predictionCol="prediction", metricName="rmse")
>>> rmse = evaluator.evaluate(predictions)
>>> new_data = [(10000, 2025, "Ford", "Mustang", "GT", "Coupe", "New", "Dealer", "New York", "NY"),
... (50000, 2038, "Toyota", "Camry", "SE", "Sedan", "Used", "Private Seller", "Los Angeles", "CA")]
>>> new_data_df = spark.createDataFrame(new_data, ["miles", "year", "make", "model", "trim", "body_type", "vehicle_type", "seller_name", "city", "state"])
>>> assembler = VectorAssembler(inputCols=["miles", "year"], outputCol="features")
>>> new_data_features = assembler.transform(new_data_df)
>>> predictions = model.transform(new_data_features)
>>> predictions.select("make", "model", "year", "miles", "prediction").show()
```

# PREDICTION

```
+-----+-----+-----+-----+-----+
| make|  model|year|miles|      prediction|
+-----+-----+-----+-----+-----+
|  Ford|Mustang|2025|10000|41127.86025576979|
|Toyota|  Camry|2038|50000|36882.95919682251|
+-----+-----+-----+-----+-----+
```

- Prediction of new data

Ford Mustang and Toyota Camry prices in the years 2025 and 2038 respectively

## BUSINESS RECOMMENDATION

- Car manufacturers should focus on producing more fuel-efficient vehicles, particularly trucks, as they are the most efficient and provide the maximum mileage while also being the least expensive option. This can help to attract more consumers who are looking for cost-effective and efficient vehicles.
- Dealers should prioritize acquiring recently bought and lightly used cars, as these tend to have a positive effect on the price, whereas mileage has a negative impact on the price. This can help dealers to maximize their profits and increase their sales.
- Companies that want to enter the market should focus on differentiating themselves by offering unique features or targeting specific niches, as the dominant players in the market are already well-established.





THANK YOU